

# SQL Programming (3)

## Readings:

- "The Database language SQL" chapter of the textbook
- SQLite tutorial <https://www.sqlitetutorial.net/>

# What we discussed in SQL1 and SQL2

- The SELECT-FROM-WHERE structure
- Single relation queries
  - The Where condition to extract rows from tables
  - Simple aggregation on whole tables
- Multi-relation queries
  - Join
  - Sub-queries

# What we discussed in SQL1 and SQL2: Exercise

Explain in English what the following SQL queries are doing.

```
select rating, length  
from movie  
where strftime('%Y', rel_date)= '2009';  
select avg(length)  
from movie;
```

Note: The "strftime()" function in SQLite is equivalent to the "to\_char()" function in SQL\*Plus. It can extract specific parts of a date value:

<https://www.sqlitetutorial.net/sqlite-date-functions/sqlite-strftime-function/>

# In this lecture

- Complex Aggregations
  - selective aggregation,
  - GROUP BY and HAVING
- SET Operators
  - UNION, INTERSECT, EXCEPT
- View
- Index

Note: EXCEPT in SQLite is equivalent to MINUS in SQL\*Plus.

# Selective Aggregation

- Aggregate the tuples selected by the WHERE clause.
- Example: What is the average length of movies produced by Roadshow?

```
SELECT AVG(length)  
FROM Movie  
WHERE studio='Roadshow';
```

| MVID | LENGTH | STUDIO   |
|------|--------|----------|
| 2    | 108    | Roadshow |
| 3    | 153    | Roadshow |




| AVG(LENGTH) |
|-------------|
| 130.5       |

# GROUP BY: Grouping Tuples

- The GROUP BY operator groups tuples -- possibly selected by a WHERE clause – into groups for aggregation.
  - Each group of tuples have unique values for the group-by attribute list.
- Example: How many movies are there for each genre?

| MVID | GENRE     |
|------|-----------|
| 1    | Drama     |
| 2    | Drama     |
| 3    | Action    |
| 3    | Adventure |
| 3    | Drama     |
| 4    | Comedy    |
| 5    | Animated  |
| 5    | Comedy    |



```
select genre, count(*)
from classification
group by genre;
```

| GENRE     | COUNT(*) |
|-----------|----------|
| Adventure | 1        |
| Animated  | 1        |
| Action    | 1        |
| Comedy    | 2        |
| Drama     | 3        |

# GROUP BY ...

- With tuples grouped into groups, groups are represented by the group-by attributes, and tuples become unrecognisable.
- With GROUP BY, only group-by attributes and aggregates for groups should be output. But SQLite does not report errors when you output non-GROUP BY attributes (although SQL\*Plus reports errors).

In SQL\*Plus:

```
1 select genre, mvID
2 from classification
3 group by genre
```

SQL> /

```
select genre, mvID
      *
```

ERROR at line 1:

ORA-00979: not a GROUP BY expression

In SQLite:

```
66 select genre, count(*), mvID
67 from classification
68 group by genre;
```

|   | genre     | count(*) | mvID |
|---|-----------|----------|------|
| 1 | Action    | 1        | 3    |
| 2 | Adventure | 1        | 3    |
| 3 | Animated  | 1        | 5    |
| 4 | Comedy    | 2        | 5    |
| 5 | Drama     | 3        | 3    |

Non-GROUP  
BY attribute.


# HAVING: select groups and aggregates

- The HAVING operator specifies conditions for choosing groups and aggregates to output.
- **Example:** What are the genres that have at least two movies? Output these genres and their total number of movies.

MVID GENRE

-----

|   |           |
|---|-----------|
| 1 | Drama     |
| 2 | Drama     |
| 3 | Action    |
| 3 | Adventure |
| 3 | Drama     |
| 4 | Comedy    |
| 5 | Animated  |
| 5 | Comedy    |



```
select genre, count(*)
from classification
group by genre
having count(*) >=2;
```

| GENRE  | COUNT(*) |
|--------|----------|
| -----  | -----    |
| Comedy | 2        |
| Drama  | 3        |




# WHERE vs. HAVING

- A WHERE clause chooses tuples for output or further aggregation. A HAVING clause chooses groups and aggregates for output.
- A WHERE clause must appear before the GROUP BY and HAVING clauses.
- **Example:** For genres starting with C or D, output those that have at least 3 movies.

MVID GENRE

-----


1 Drama  
2 Drama  
3 Action  
3 Adventure  
3 Drama  
4 Comedy  
5 Animated  
5 Comedy




```
select genre, count(*)
from classification
where genre like 'C%' or genre like 'D%'
group by genre
having count(*) >=3
```

| GENRE | COUNT(*) |
|-------|----------|
| ----- |          |
| Drama | 3        |


# Set operations



A union B --  
elements in A or B



A minus B -- elements in  
A but not B



A intersect B --  
elements in A and B

A and B are sets of elements of the same structure/domain.

- If A contains integers B must also contains integers.
- If A is a set of tuples of 3 components, then so is B.

# Set operators:

## Union, Intersect and Minus

- The Union, Intersect and Except operators applied to relations are expressed by the following expressions possibly involving subqueries:
  - <subquery> UNION <subquery>
  - <subquery> INTERSECT <subquery>
  - <subquery> EXCEPT <subquery>

# Our previous failed query

Find the movies (mvID) that have both “Marie Gillain” and “Audrey Tautou”.

```
select mvID  
from Cast  
where actor='Marie Gillain'  
       and actor='Audrey Tautou';
```

# INTERSECT

- Solution:
  1. Find the movies that have “Marie Gillain”.
  2. Find the movies that have “Audrey Tautou”.
  3. Find the intersection of results from 1 and 2, which is the solution.

# Solution

Movies that have  
"Marie Gillain".



```
Select mvID  
From Cast  
Where actor='Marie Gillain'
```

INTERSECT

```
Select mvID  
From Cast  
Where actor='Audrey Tautou';
```

Movies that have  
"Audrey Tautou".

# UNION

- Find the movies (mvID) that have either “Tom Hanks” or “Audrey Tautou”.
- Solution:
  1. Find the movies by “Tom Hanks”.
  2. Find the movies by “Audrey Tautou”.
  3. Union the results from 1 and 2, and form the solution.

# Solution

```
select mvID
from Cast
where actor='Tom Hanks'
UNION
select mvID
from Cast
where actor='Audrey Tautou';
```

| MVID  |
|-------|
| ----- |
| 1     |

| MVID  |
|-------|
| ----- |
| 2     |

| MVID  |
|-------|
| ----- |
| 1     |
| 2     |



# EXCEPT

Find the movies that are *only* in the genre "Drama".

1. Find the movies in the genre "Drama".
2. Find the movies in genres other than "Drama".
3. Take the difference of 1 and 2, and form the solution.

Note that EXCEPT in SQLite is equivalent to MINUS in SQL\*PLUS.

# Solution

```
select mvID
from Classification
where genre='Drama'
except
select mvID
from Classification
where genre != 'Drama';
```

| MVID  |
|-------|
| ----- |
| 1     |
| 2     |

| MVID  |
|-------|
| ----- |
| 1     |
| 2     |
| 3     |

| MVID  |
|-------|
| ----- |
| 3     |
| 3     |
| 4     |
| 5     |
| 5     |

# Set operators Remove duplicates!

- The default SELECT-FROM-WHERE statement keeps duplicates (the bag semantics).
- The default union, intersection, and except expressions remove duplicates (set semantics).

# Set operators remove duplicates ...

- Find the movies (mvID) that are not in the genre "Drama"

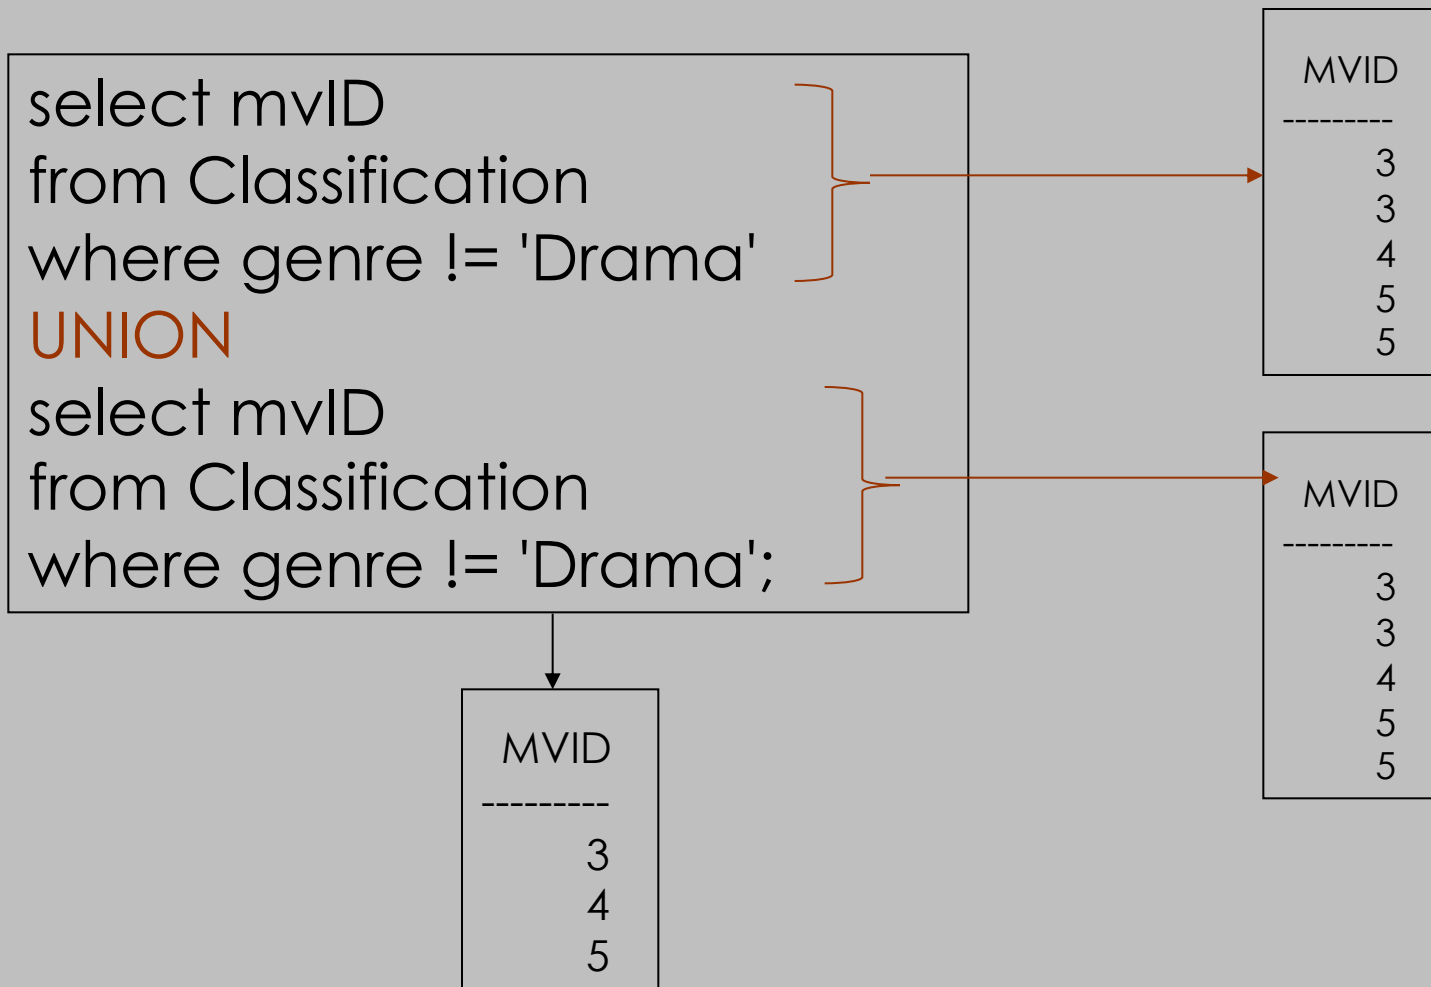
```
select mvID  
from Classification  
where genre != 'Drama';
```

| MVID |
|------|
| 3    |
| 3    |
| 4    |
| 5    |
| 5    |

```
select DISTINCT mvID  
from Classification  
where genre != 'Drama';
```

| MVID |
|------|
| 3    |
| 4    |
| 5    |

# Set operators remove duplicates ...



# Views

- A view is a virtual table defined by an SQL query.
- Views do not keep data, but can be queried. When a view is queried, it is replaced by its definition to execute the query.
- Views generally are for not for update.
- View definition:

```
CREATE VIEW <name>AS <view-definition>
```

# Views ...

- Define a view for the total number of actors for each movie – attributes can be renamed.
- The view can be queried and queries often become simpler.

```
CREATE VIEW genreCount (mvid,numGenre) AS  
SELECT mvid, count(genre)  
FROM Classification  
group by mvid;
```

Querying the view GenreCount(mvid, numGenre) :

```
select mvid  
from GenreCount  
where numGenre >1;
```

```
select avg(numGenre)  
from GenreCount;
```

# Indexes

- CS convention: **indexes** not indices.
- An **Index** is a data structure used to speed up access to tuples of a relation.
  - A DBMS uses an index on a table to search for a row rather than scanning the whole table. This greatly reduces search time and disk input/output.
- An index for a table is always a balanced search tree with giant nodes (a full disk page) called a B-tree.
  - This topic will be discussed in more details in the course *Database Systems (COSC2406/2407)*.



## Creating Indexes

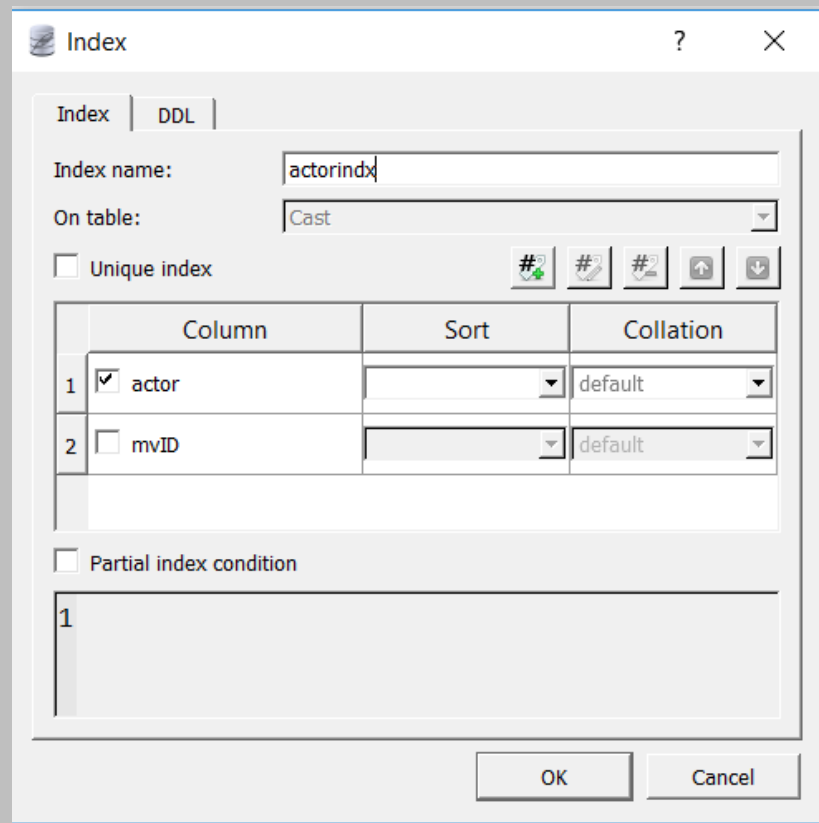
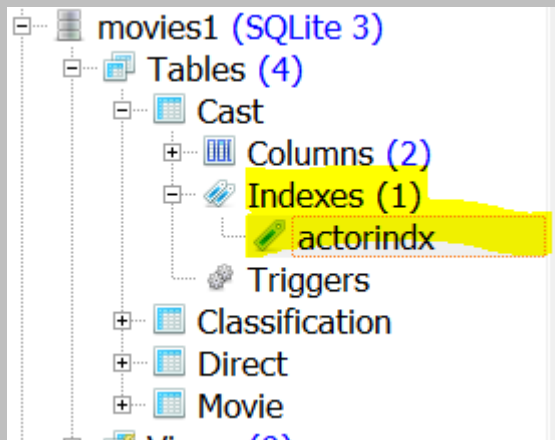
- In most if not all DBMSs, implicit indexes are created automatically when the PRIMARY KEY constraint is defined.
- An index can also be created in SQLite using a CREATE INDEX statement.

```
CREATE INDEX indexname  
    ON tablename (col1, col2, ...);
```

# Creating Indexes ...

- An index named actorindx is created.

CREATE INDEX actorindx  
ON Cast(actor);



## Using Indexes

- When an index is created, a user does not need to open or use the index with a command. Rather an index is used by the DBMS for processing queries.
- Indexes are kept separately from their base tables.
- Every insertion or deletion in a table updates the index, which is added overhead on the system.

# Using Index ...

```
CREATE INDEX actorindx  
ON Cast(actor);
```

The DBMS SQL engine will automatically use the “actorindx” index when processing queries involving the column “actor”. Queries run more efficiently.

```
select *  
from Cast  
where actor > 'A%';
```

```
select C1.mvid, C2.mvid, C1.actor  
from Cast C1, Cast C2  
where C1.actor=C2.actor  
and C1.mvid < C2.mvid;
```

# Database Tuning

- A main task in making a database run fast is deciding which indexes to create.
  - An index speeds up queries that can use it.
  - An index slows down all modifications on its relation because the index must be modified too.

## Database Tuning ...

- Generally an index is created on a column if the column
  - is used very often in querying and joining,
  - has a big domain of values, or
  - contains many Null values.
    - Null values are removed in indexes.
- An index should not be created for
  - a very small table,
  - a column not used often in queries, or
  - a table that often gets updated.